

Midterm Exam I- Solutions

ECE 685D– Introduction to Deep Learning

Fall 2023

Instructor: Prof. Vahid Tarokh
ECE Department, Duke University

Oct 4th 2023
10:05 AM - 11:20 AM
(Exam duration: 75 minutes)

Name: _____
Duke ID: _____

This exam contains 10 pages and 9 questions. This exam has 105 points of which 5 are bonus points. This is a closed-book exam. No exam aids are allowed. You are not allowed to communicate with others.

Distribution of Marks

Question	Points	Score
1	2	
2	2	
3	3	
4	3	
5	10	
6	5	
7	15	
8	30	
9	35	
Total:	105	

For each of the following questions, circle the letter of your choice. Each question has AT LEAST one correct option unless explicitly mentioned.

1. (2 points) In a logistic regression model, the decision boundary can be ...

- (a) linear
- (b) non-linear
- (c) both (a) and (b)

Answer: (c)

2. (2 points) Among these commonly-used CNN layers, what is the least computationally complex in terms of floating point operations?

- (a) Conv layer (convolution operation + bias addition)
- (b) Average pooling
- (c) Max pooling
- (d) Batch Normalization

Answer: (c)

3. (3 points) Which of the following optimization methods uses first-order momentum?

- (a) RMSProp
- (b) Gauss-Newton
- (c) Adam
- (d) Stochastic Gradient Descent

Answer: (c)

4. (3 points) Making your network deeper by adding more parametrized layers will always...

- (a) reduce the training loss.
- (b) improve the performance on unseen data.
- (c) slow down training and inference speed.
- (d) both (a) and (b)

Answer: (c)

5. (10 points) Please complete the Pytorch code below by providing the missing lines necessary to train the CNN model (do not need to worry about syntax)

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.utils.data.DataLoader as dataloader

# Load the data
train_data = torch.load("train_data.pt")
trainloader = dataloader(train_data)
test_data = torch.load("test_data.pt")
trainloader = dataloader(test_data)

# Define the model
model = nn.Sequential(
    nn.Conv2d(3, 64, kernel_size=3),
    nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Conv2d(64, 128, kernel_size=3),
    nn.ReLU(),
    nn.MaxPool2d(2),
    nn.Flatten(),
    nn.Linear(128 * 7 * 7, 1000),
    nn.ReLU(),
    nn.Linear(1000, 10),
)

# Train the model
# Write in your code here:
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

for epoch in range(30):
    for i, data in enumerate(trainloader):
        input, labels = data
        optimizer.zero_grad()
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

# Evaluate the model
accuracy = model.evaluate(test_data)
print("Accuracy:", accuracy)
```

6. (5 points) A function $f(\cdot) : \mathbb{R} \rightarrow \mathbb{R}$ is said to be piece-wise constant if for some non-negative integer n if there exists increasing real numbers

$$-\infty = a_0 < a_1 < a_2 < \cdots < a_n < a_{n+1} = \infty$$

and real numbers c_0, c_1, \dots, c_n such that $f(x) = c_j$ for $a_j < x \leq a_{j+1}$ for $j = 0, 1, \dots, n-1$ and $f(x) = c_n$ for $a_n < x$.

Can a piece-wise constant function be suitable for use as an activation function in a neural network? Please explain your answer.

Answer: No, since it is mostly non-differential. Hence, it cannot backpropagate weights to previous layers.

7. Consider a binary logistic regression problem as follows:

$$p_i = p(y_i = 1 | \mathbf{x}_i) = \sigma(\mathbf{w}^T \mathbf{x}_i + b), \forall i \in \{1, 2\} \quad (1)$$

where $y \in \{1, 0\}$, $\mathbf{x} \in \mathbb{R}^{2 \times 1}$, $\mathbf{w} \in \mathbb{R}^{2 \times 1}$, $b \in \mathbb{R}$, and $\sigma(\cdot)$ is the sigmoid function, given as: $\sigma(a) = 1/(1 + e^{-a})$. Given a dataset with two data points $\{\mathbf{x}_1, y_1\} = \{(1, 0)^T, 2\}$, $\{\mathbf{x}_2, y_2\} = \{(1, 1)^T, 3\}$.

The loss function is

$$\mathcal{L} = - \sum_{i=1}^2 y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \quad (2)$$

The initial value of \mathbf{w} is $\mathbf{w}_1 = (1, 1)^T$.

The initial value of b is $b_1 = 0.5$.

- (10 points) Write out the gradient of the loss function \mathcal{L} with respect to the weight \mathbf{w} and bias b explicitly.
- (5 points) Perform one step of Nesterov's accelerated gradient descent method with $\beta = 0.5$ on \mathbf{w} and b using the dataset formed with two data points $\{\mathbf{x}_i, y_i\}_{i=1}^2$. **Use the definition of Nesterov's Accelerated Gradient Descent in the lecture notes, as shown below** to calculate t_2, \mathbf{w}_2, b_2 .

Algorithm 1 Nesterov's Accelerated Gradient Descent

First define the following sequences: $\lambda_0 = 0$, $\lambda_k = (1 + \sqrt{1 + 4\lambda_{k-1}^2})/2$, $\gamma_k = (1 - \lambda_k)/\lambda_{k+1}$
for $k = 1, 2, \dots$ **do**
 $\mathbf{t}_{k+1} = \mathbf{w}_k - \nabla \mathcal{L}(\mathbf{w}_k)/\beta$
 $\mathbf{w}_{k+1} = (1 - \gamma_k)\mathbf{t}_{k+1} + \gamma_k \mathbf{t}_k$
end for

(a)

$$\begin{aligned} \frac{\partial L}{\partial w} &= - \sum_{i=1}^2 \frac{y_i}{p_i} \frac{\partial p_i}{\partial w} + \frac{1 - y_i}{1 - p_i} \frac{\partial (1 - p_i)}{\partial w} \\ &= - \sum_{i=1}^2 \frac{y_i}{p_i} p_i (1 - p_i) x_i - \frac{1 - y_i}{1 - p_i} p_i (1 - p_i) x_i \\ &= - \sum_{i=1}^2 (y_i - p_i) x_i \end{aligned}$$

$$\begin{aligned}
\frac{\partial L}{\partial b} &= - \sum_{i=1}^2 \frac{y_i}{p_i} \frac{\partial p_i}{\partial b} + \frac{1 - y_i}{1 - p_i} \frac{\partial (1 - p_i)}{\partial b} \\
&= - \sum_{i=1}^2 \frac{y_i}{p_i} p_i (1 - p_i) - \frac{1 - y_i}{1 - p_i} p_i (1 - p_i) \\
&= - \sum_{i=1}^2 (y_i - p_i)
\end{aligned}$$

(b) $\lambda_0 = 0, \lambda_1 = 1, \gamma_1 = 0$

We have: $\nabla L(w, b) = \left(\frac{\partial L}{\partial w}, \frac{\partial L}{\partial b} \right)$

From Nesterove Algorithm, we obtain: $t_2 = (w, b) - \nabla L(w, b) / \beta$

Hence: $w_2 = (1 - \gamma_1)t_2[0] + \gamma_1 t_1[0] = t_2[0]$ and $b_2 = (1 - \gamma_1)t_2[1] + \gamma_1 t_1[1] = t_2[1]$

8. Consider an RGB image $X = [X_0, X_1, X_2]$ with three channels, and given as follows:

$$X_0 = \begin{bmatrix} 2 & 1 & 0 & 0 \\ 0 & 0 & 2 & 1 \\ 0 & 2 & 0 & 1 \\ 2 & 1 & 0 & 1 \end{bmatrix}, X_1 = \begin{bmatrix} 2 & 2 & 0 & 0 \\ 0 & 0 & 2 & 1 \\ 0 & 0 & 2 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix}, X_2 = \begin{bmatrix} 2 & 1 & 0 & 0 \\ 0 & 0 & 2 & 1 \\ 2 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix}. \quad (3)$$

This image is passed through the convolutional filter with the weights $W = [W_0, W_1, W_2]$ of size $3 \times 3 \times 3$, step size of 1, and is given as follows:

$$W_0 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -2 & 0 \\ 0 & 0 & -1 \end{bmatrix}, W_1 = \begin{bmatrix} 1 & 2 & 0 \\ 2 & 0 & -1 \\ 0 & -1 & 1 \end{bmatrix}, W_2 = \begin{bmatrix} 0 & 0 & -2 \\ 0 & 1 & 2 \\ -2 & 2 & 0 \end{bmatrix}. \quad (4)$$

The output of the convolutional filter is given as follows:

$$Y = ReLU \left(\sum_{i=0}^2 (X'_i * W_i) + 2 \times 1_{4 \times 4} \right) \quad (5)$$

where Y is the output image, X' is the input image after applying zero-padding around the edges (i.e. each channel is converted to a 6×6 matrix such that a row of zeros is added to the top and bottom and a column of zeros is added to the left and right.), $X'_i * W_i$ is the convolution of the i -th channel of X' with the i -th channel of W , and $1_{4 \times 4}$ is a 4×4 matrix with all ones.

- (a) (20 points) Compute the output Y of the image X .
- (b) (5 points) Apply max pooling on non-overlapping 2×2 sub-matrices of the output image and compute the output.
- (c) (5 points) Apply average pooling on non-overlapping 2×2 sub-matrices of the output image and compute the output.

Answer:

- (a) Here are the results for the non-flipped kernel. Note that we also give full credit for the flipped kernel approach.

$$X_0 * W_0 = \begin{bmatrix} -4 & -4 & -1 & 0 \\ -2 & 2 & -4 & -2 \\ -1 & -4 & -1 & 0 \\ -4 & -2 & 2 & -2 \end{bmatrix} \quad (6)$$

$$X_1 * W_1 = \begin{bmatrix} -2 & 6 & 3 & -1 \\ 4 & 6 & -1 & 4 \\ 1 & -3 & 5 & 7 \\ -1 & 0 & 5 & 2 \end{bmatrix} \quad (7)$$

$$X_2 * W_2 = \begin{bmatrix} 4 & 1 & 4 & -2 \\ 2 & 0 & 4 & 1 \\ 2 & -2 & -4 & 2 \\ 2 & 1 & 2 & 1 \end{bmatrix}. \quad (8)$$

$$Y = ReLU \left(\sum_{i=0}^2 (X'_i * W_i) + 2 \times 1_{4 \times 4} \right) = \begin{bmatrix} 0 & 5 & 8 & 0 \\ 6 & 10 & 1 & 5 \\ 4 & 0 & 2 & 11 \\ 0 & 1 & 11 & 3 \end{bmatrix} \quad (9)$$

(b)

$$maxpool(Y) = \begin{bmatrix} 10 & 8 \\ 4 & 11 \end{bmatrix} \quad (10)$$

(c)


$$avepool(Y) = \begin{bmatrix} 21/4 & 7/2 \\ 5/4 & 27/4 \end{bmatrix} \quad (11)$$

9. Given the following neural network with two input units (i_1, i_2) , fully-connected layers and ReLU activations. The weights and bias of hidden units are denoted w and b , with h_1, h_2, h_3, h_4 are ReLU units.

$$h_1 = \text{ReLU}(i_1 w_{11} + i_2 w_{21} + b_1) \quad (12)$$

The outputs are denoted as (y_1, y_2) , and the ground truth targets are denoted as (t_1, t_2) .

$$y_1 = \text{ReLU}(h_1 w_{31} + h_2 w_{41} + b_3) \quad (13)$$



The values of the variables are given as follows:

i_1	i_2	w_{11}	w_{12}	w_{21}	w_{22}	w_{31}	w_{32}	w_{41}	w_{42}	b_1	b_2	b_3	b_4	t_1	t_2
1	2	1	0.5	-0.5	1	0.5	-2	-1	0.5	-0.5	-0.5	1	1	2	4

- (10 points) Compute the output (y_1, y_2) of the input (i_1, i_2) using the network parameters as specified above (please write down all calculations of the intermediate layers)
- (5 points) Compute the mean squared error of the computed output (y_1, y_2) and the target labels (t_1, t_2) .
- (5 points) Using the calculated MSE above, update the weight w_{31} using gradient descent and backpropagation algorithm with a learning rate of 0.01 (write down all your computations).
- (5 points) Using the calculated MSE above, update the weight w_{42} using gradient descent and backpropagation algorithm with a learning rate of 0.01 (write down all your computations).
- (10 points) Using the calculated MSE above, update the weight w_{22} using gradient descent and backpropagation algorithm with a learning rate of 0.01 (write down all your computations).

Answer:

(a) First, $h_1 = 0, h_2 = 2$.

Next, $y_1 = h_3 = \text{ReLU}(h_1 w_{31} + h_2 w_{41} + b_3) = 0, y_2 = 2$.

(b) $L_{MSE} = \frac{1}{2} \sum_{i=1}^2 (y_i - t_i)^2 = 4$

(c) $\frac{\partial L}{\partial w_{31}} = \frac{\frac{1}{2} \sum_{i=1}^2 (y_i - t_i)^2}{\partial w_{31}} = (y_1 - t_1) \frac{\partial h_3}{\partial w_{31}}$.

Since input of $h_3 < 0$, $\frac{\partial h_3}{\partial w_{31}} = 0$.

Hence, $w_{31} = w_{31} - 0.01 \frac{\partial L}{\partial w_{31}} = 0.5 - 0 = 0.5$.

(d) $\frac{\partial L}{\partial w_{42}} = \frac{\frac{1}{2} \sum_{i=1}^2 (y_i - t_i)^2}{\partial w_{42}} = (y_2 - t_2) \frac{\partial h_4}{\partial w_{42}} = (y_2 - t_2) h_2 = -4$.

Hence, $w_{42} = w_{42} - 0.01 \frac{\partial L}{\partial w_{42}} = 0.54$.

(e) $\frac{\partial L}{\partial w_{22}} = (y_1 - t_1) \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial w_{22}} + (y_2 - t_2) \frac{\partial h_4}{\partial h_2} \frac{\partial h_2}{\partial w_{22}} = -2$.

Hence, $w_{22} = w_{22} - 0.01 \frac{\partial L}{\partial w_{22}} = 1.02$.